

CSE 124: NETWORKING FUNDAMENTALS

George Porter
Sept 27, 2022



ATTRIBUTION

- These slides are released under an Attribution-NonCommercial-ShareAlike 3.0 Unported (CC BY-NC-SA 3.0) Creative Commons license

AGENDA

Today:

- Basics of networking, addressing, performance metrics

Tuesday:

- Basics of Go, naming and addressing

TODO THIS WEEK

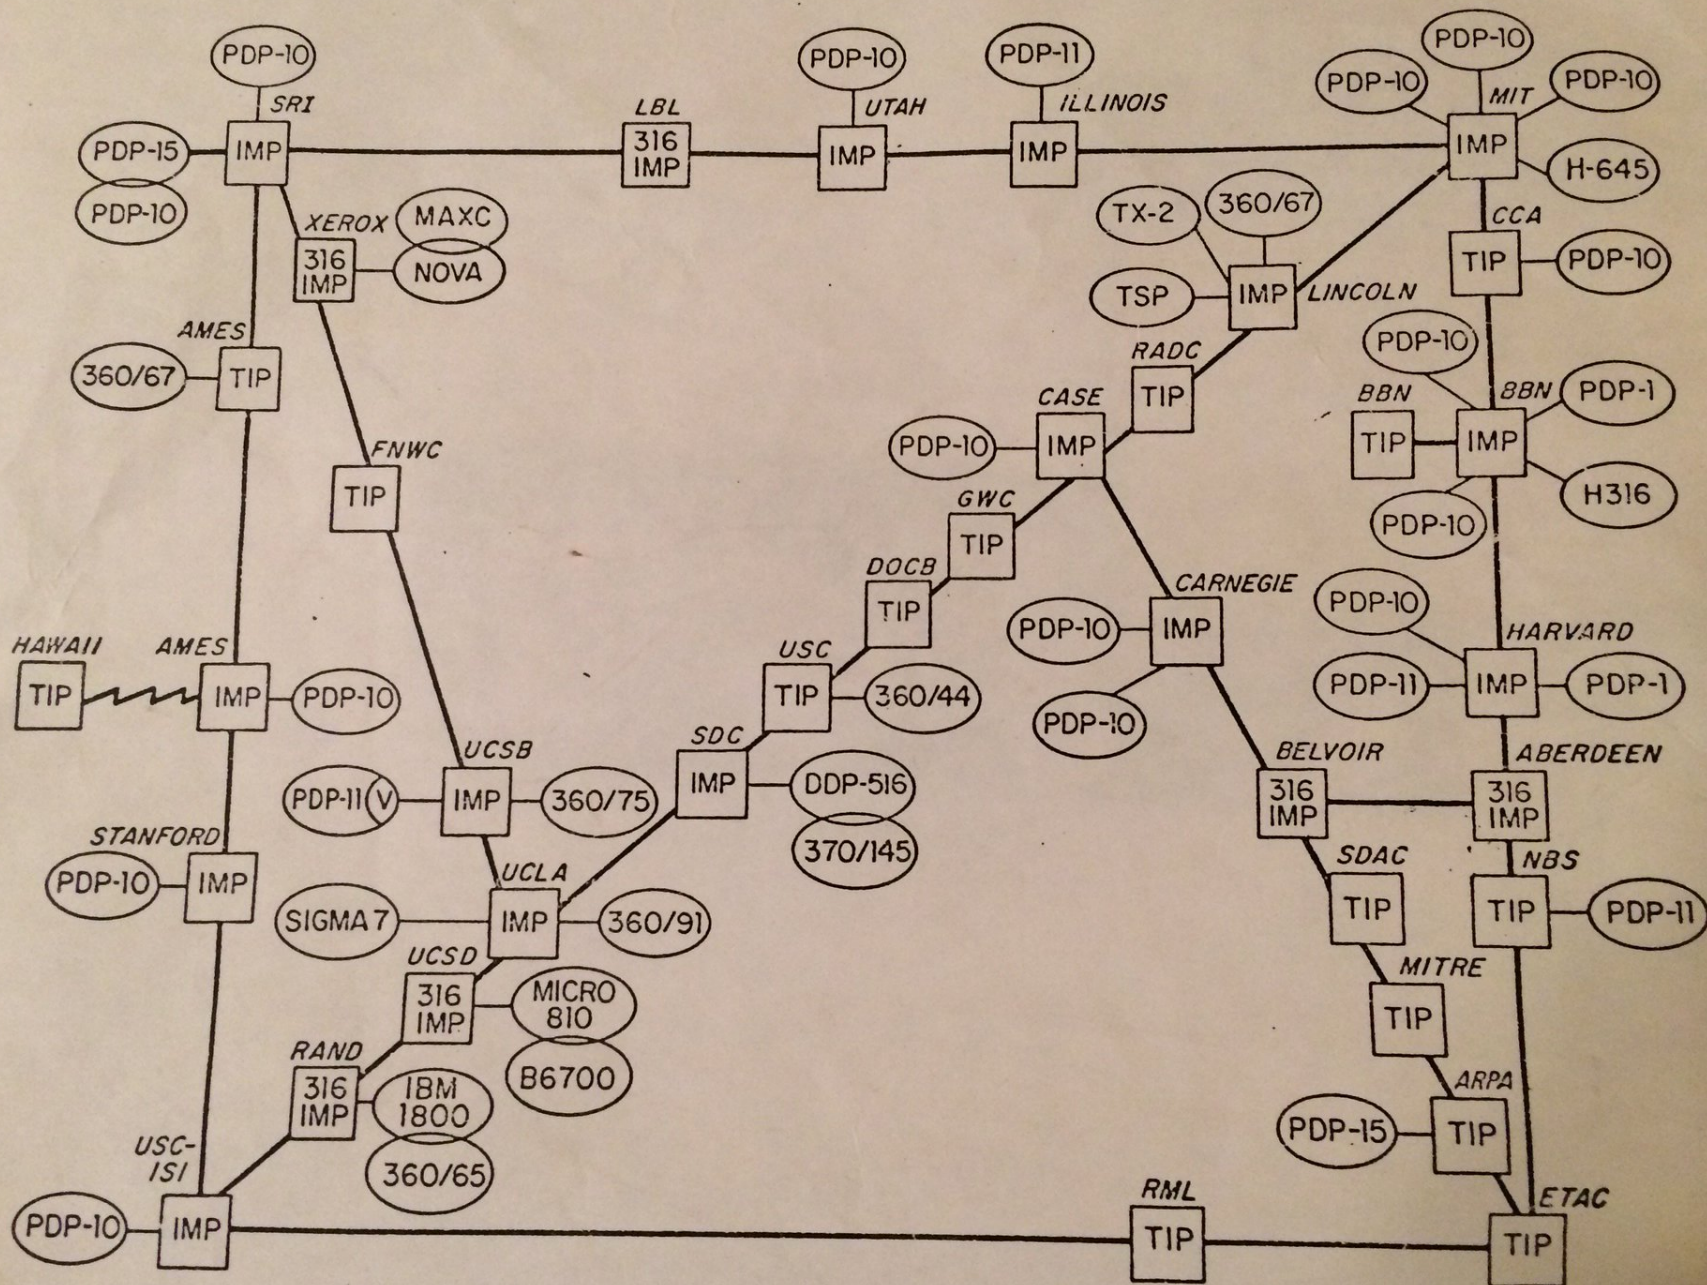
1. Read Network Programming with Go chapters 1, 2, and 3
2. Project 1 is now out (due Fri Sept 30 at 5pm)

BRIEF HISTORY OF THE INTERNET

- 1968 - DARPA (Defense Advanced Research Projects Agency) contracts with BBN (Bolt, Beranek & Newman) to create ARPAnet
- 1970 - First five nodes:
 - UCLA
 - Stanford
 - UC Santa Barbara
 - U of Utah, and
 - BBN
- 1974 - TCP specification by Vint Cerf
- 1984 – On January 1, the Internet with its 1000 hosts converts en masse to using TCP/IP for its messaging

Data from the Internet Society

ARPA NETWORK, LOGICAL MAP, MAY 1973





Outline

1. Performance
2. Layering
3. Addressing

PERFORMANCE METRICS



- Bandwidth: number of bits transmitted per unit of time
- Latency = Propagation + Transmit + Queue
 - Propagation = Distance/SpeedOfLight(*)
 - Transmit = 1 bit/Bandwidth
 - Queue = Time waiting in switches/routers behind other traffic (traffic jam)
- Overhead
 - # secs for CPU to put message on wire
- Error rate
 - Probability P that message will not arrive intact

* In that particular medium

BANDWIDTH VS. LATENCY

1 Byte Object

	Latency: 1 ms	Latency: 100 ms
Bandwidth: 1 Mbps	1,008 μ s	100,008 μ s
Bandwidth: 100 Mbps	1,000 μ s	100,000 μ s

10 MB Object

	Latency: 1 ms	Latency: 100 ms
Bandwidth: 1 Mbps	80.001 s	80.1 s
Bandwidth: 100 Mbps	.801 s	.9 s

NETWORK PERFORMANCE MEASUREMENT UNITS

Exp.	Explicit	Prefix	Exp.	Explicit	Prefix
10^{-3}	0.001	milli	10^3	1,000	Kilo
10^{-6}	0.000001	micro	10^6	1,000,000	Mega
10^{-9}	0.000000001	nano	10^9	1,000,000,000	Giga
10^{-12}	0.0000000000001	pico	10^{12}	1,000,000,000,000	Tera
10^{-15}	0.0000000000000001	femto	10^{15}	1,000,000,000,000,000	Peta
10^{-18}	0.0000000000000000001	atto	10^{18}	1,000,000,000,000,000,000	Exa
10^{-21}	0.0000000000000000000001	zepto	10^{21}	1,000,000,000,000,000,000,000	Zetta
10^{-24}	0.000000000000000000000001	yocto	10^{24}	1,000,000,000,000,000,000,000,000	Yotta

TERMINOLOGY STYLE

- Mega versus Meg, Kilo versus Kilo
 - Computer architecture: Mega $\rightarrow 2^{20}$, Kilo $\rightarrow 2^{10}$
 - Computer networks: Mega $\rightarrow 10^6$, Kilo $\rightarrow 10^3$
- Mbps versus MBps
 - Networks: typically megabits per second
 - Architecture: typically megabytes per second
- Bandwidth versus throughput
 - Bandwidth: available over link
 - Throughput: available to application
 - E.g. subtract protocol headers, etc.

PERFORMANCE TOOLS

- Ping
 - Test if other side is “alive”
 - Measures round-trip latency
- iperf3
 - Times how long it takes to send N bytes to the other endpoint
 - Used to calculate bandwidth



Outline

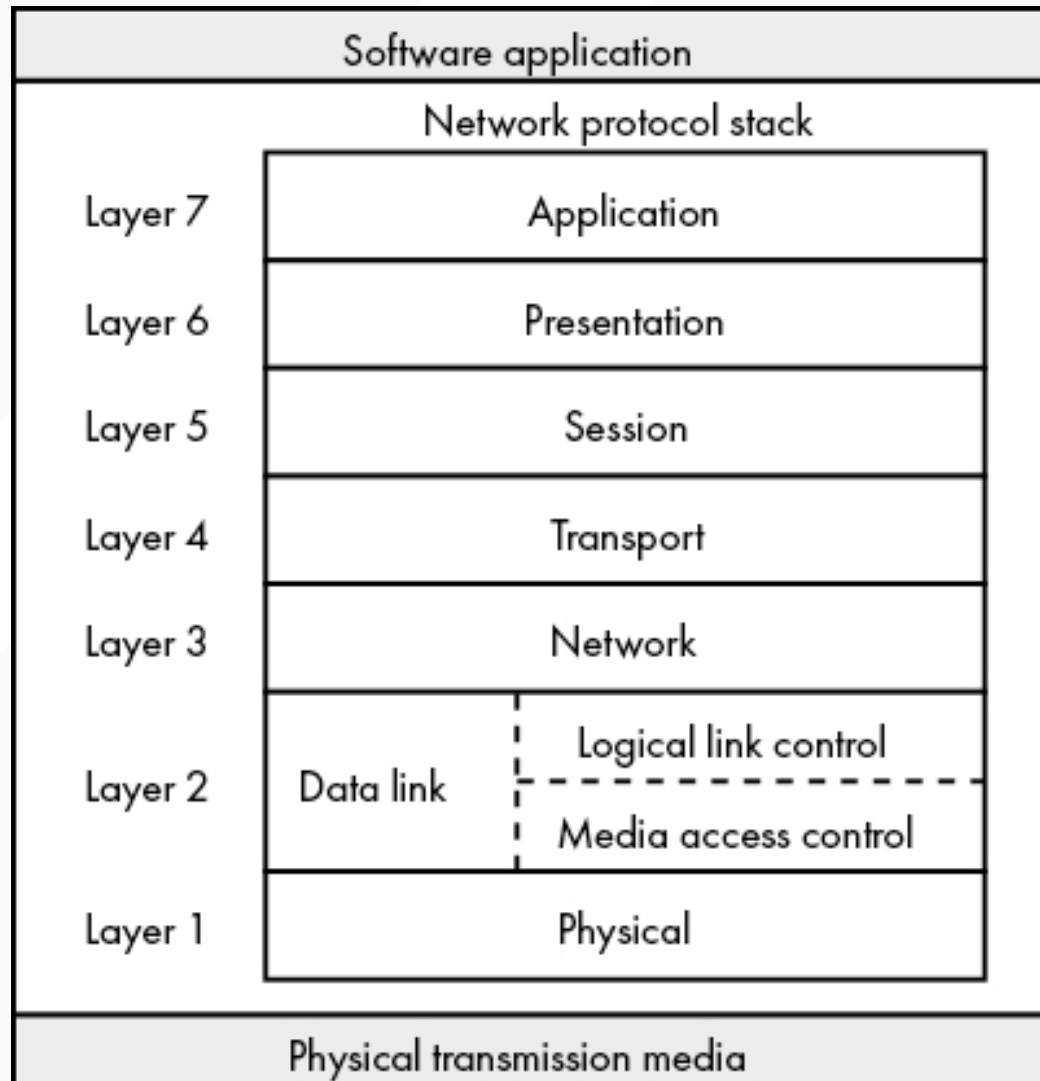
1. ~~Performance~~
2. Layering
3. Addressing

NETWORKING AND LAYERING

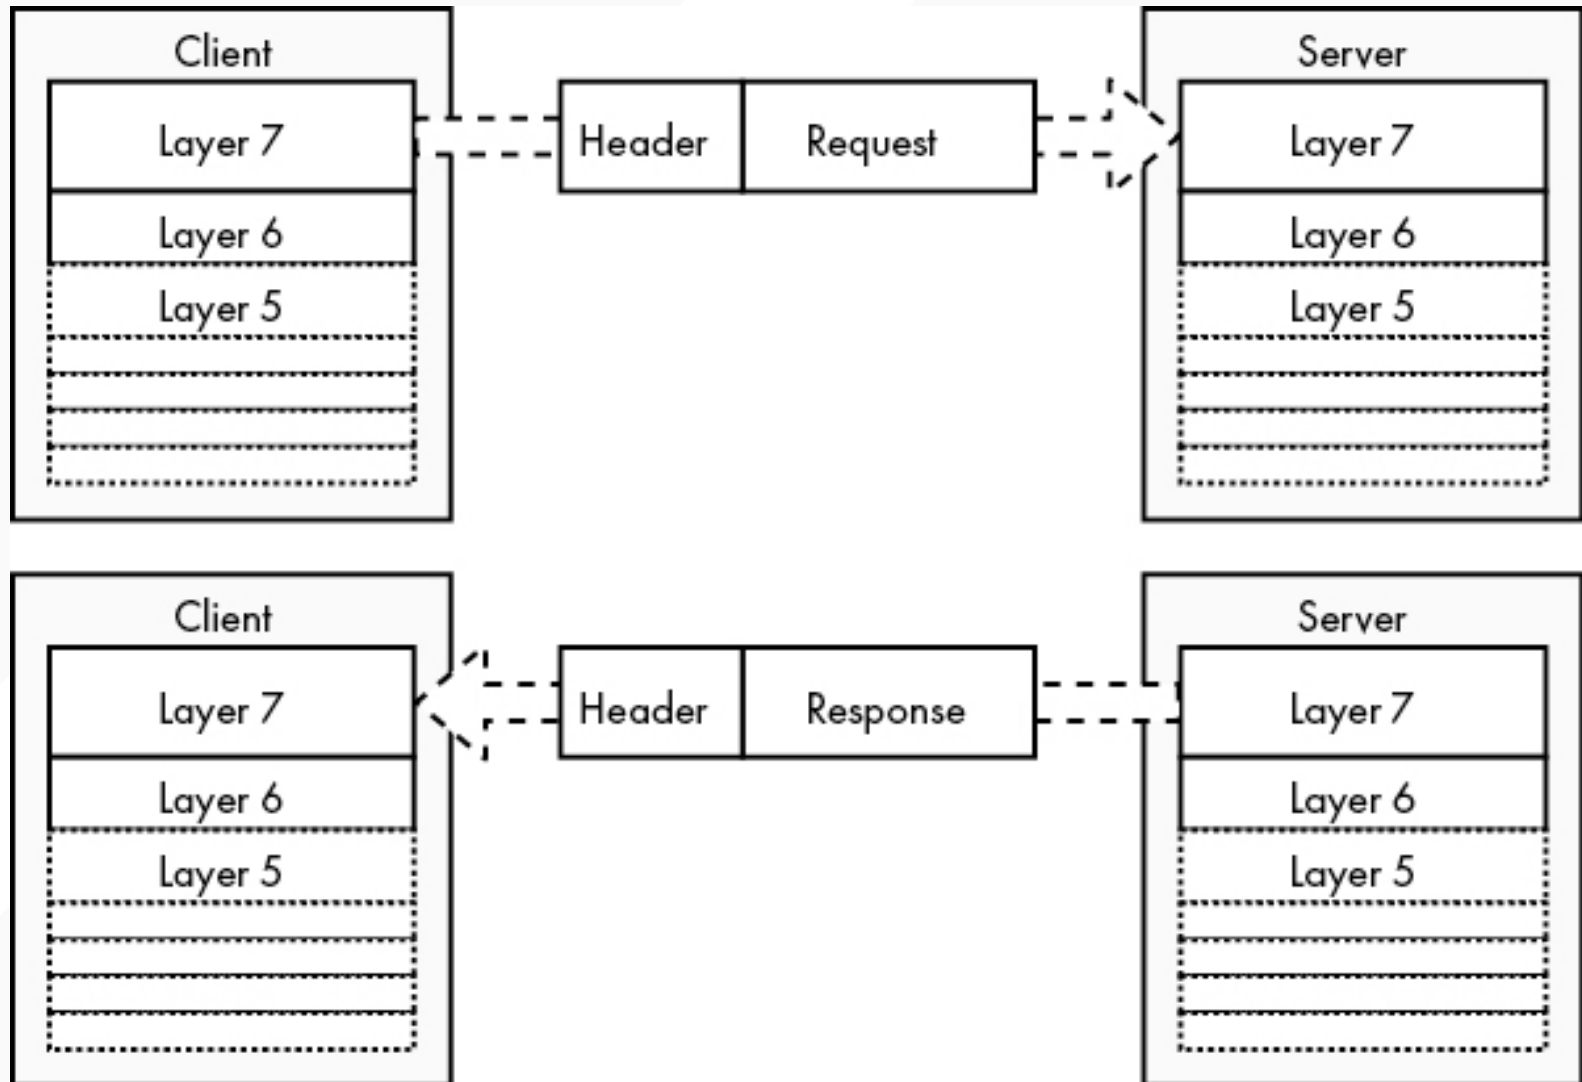
“All problems in computer science can be solved by another level of indirection.”

David J. Wheeler

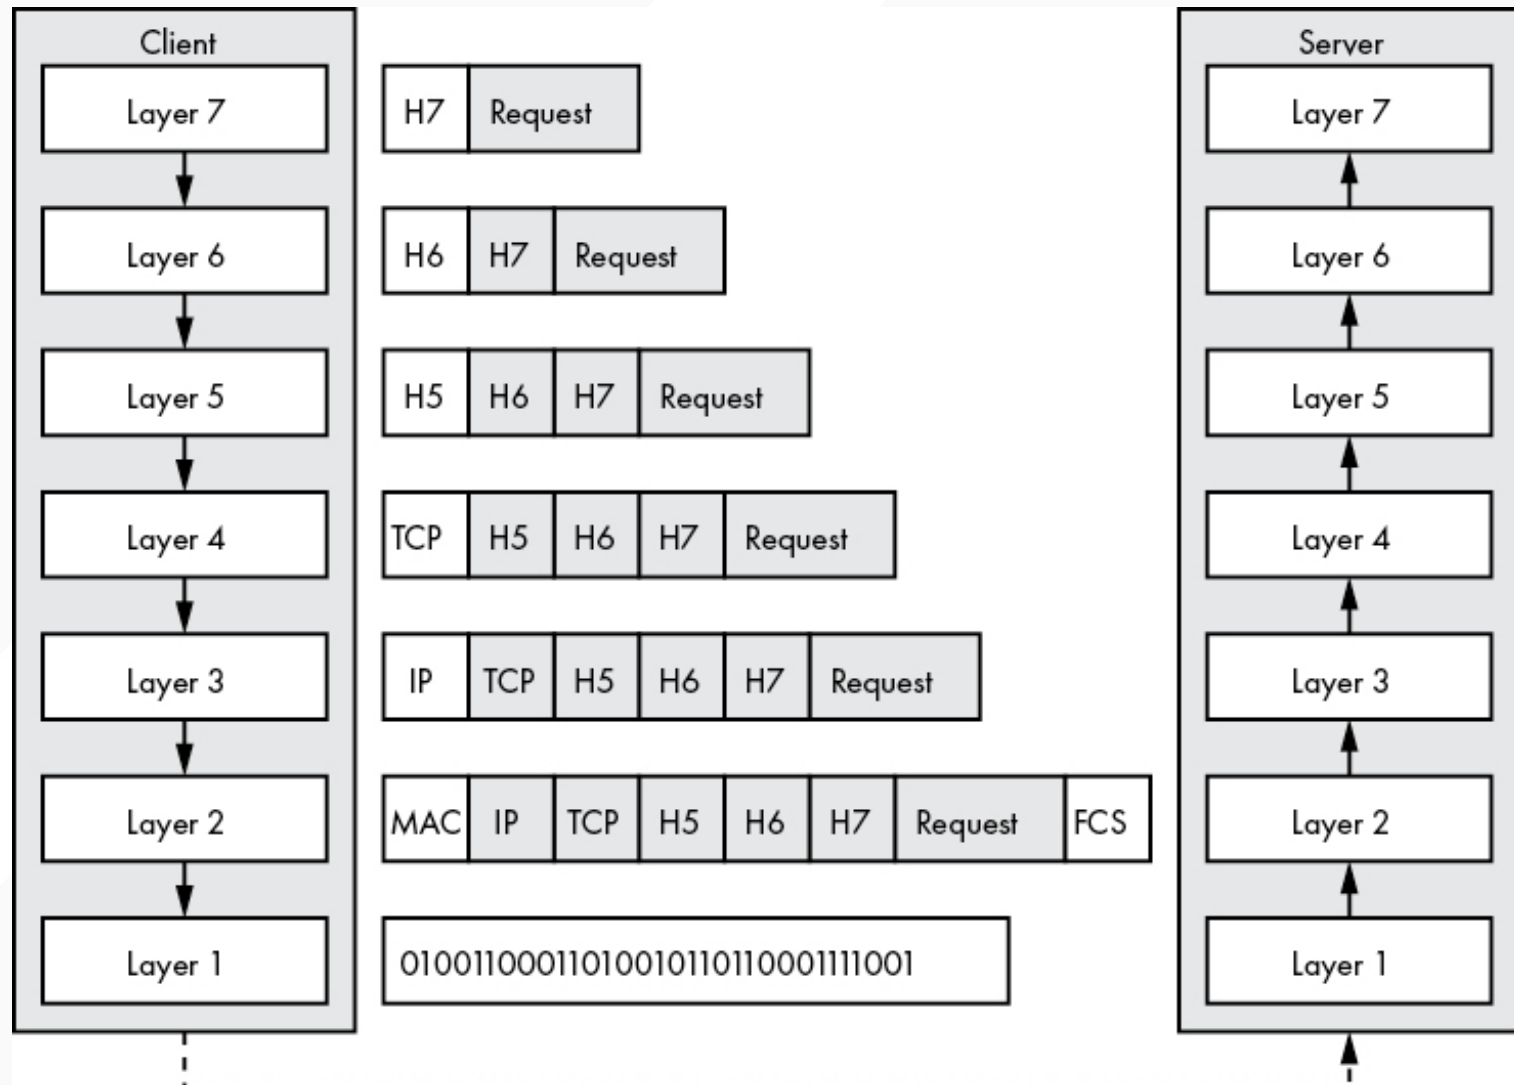
OSI NETWORK STACK



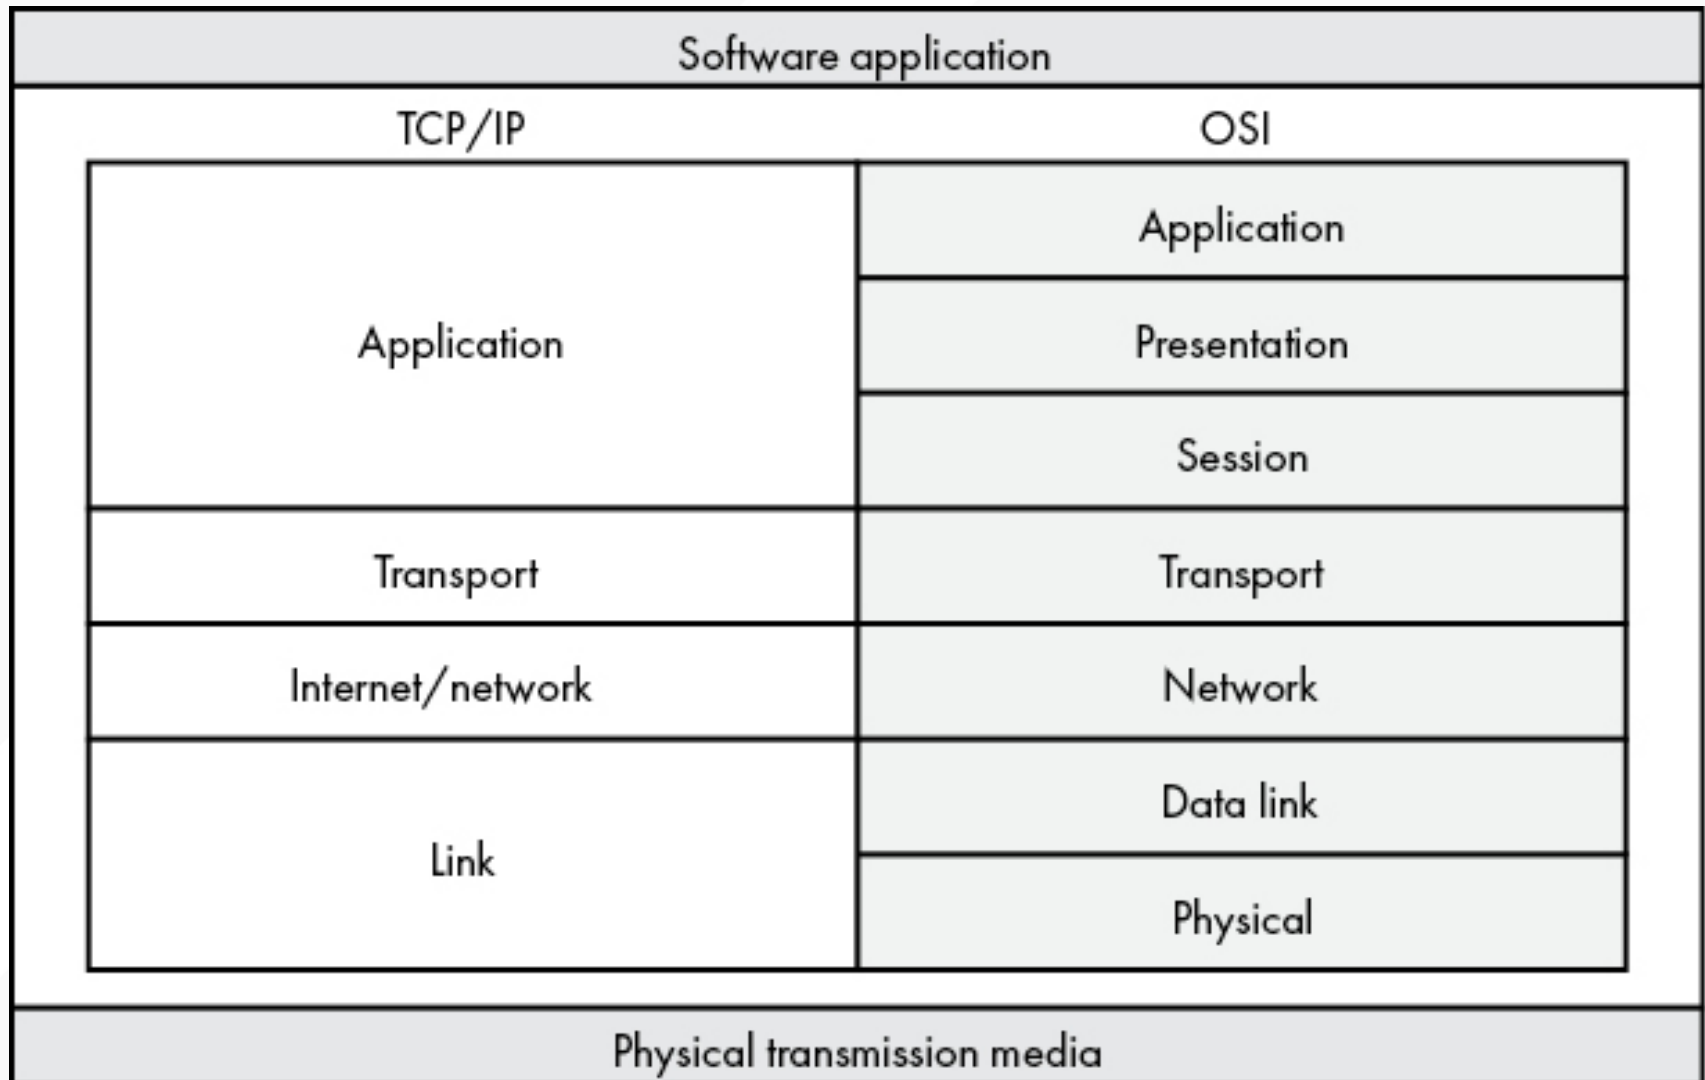
PROTOCOLS GOVERN MESSAGES EXCHANGED WITHIN A SINGLE LAYER



LAYERING AND ENCAPSULATION



TCP/IP MODEL (VS OSI MODEL)

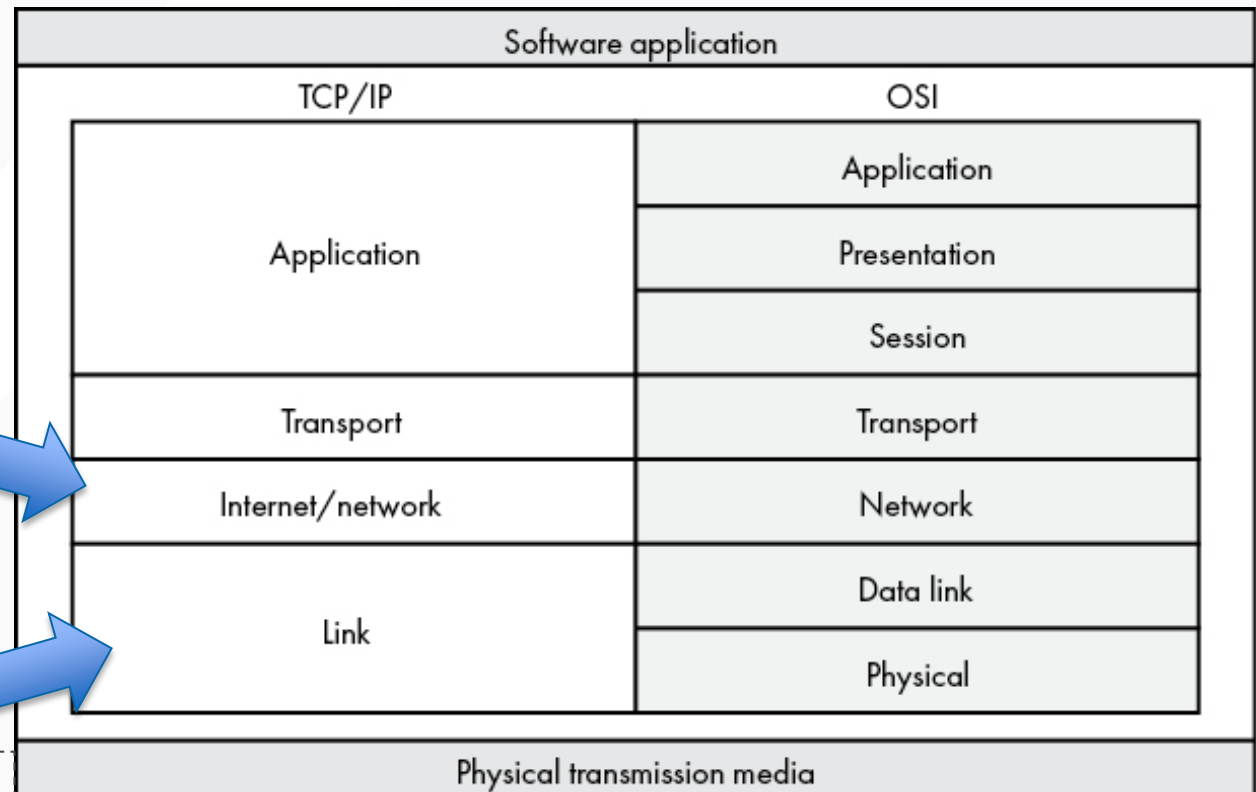


TCP/IP MODEL (VS OSI MODEL)

Forwarding packets
between subnetworks
(E.g. IPv4 or IPv6)



Forwarding packets
(frames) within one
subnetwork
(E.g. Ethernet)





Outline

1. ~~Performance~~
2. ~~Layering~~
3. Addressing

MOTIVATION FOR OUR DISCUSSION OF ADDRESSING

1. When you implement a *client* application, you will typically need to communicate with a remote system/service via its *IP Address*
2. When you implement a cloud-hosted *network service*, you will typically be assigned a *subnet* of IP addresses to use for your various server programs, and will need to use and manage those appropriately

IP PROTOCOL

- Scenario: An internet connected device wants to send a message to another internet connected device anywhere in the world
 - IP protocol handles this
 - Prepend an “IP header” to the message, set the destination IP address, set the source IP address, and Internet routers will forward it along the shortest path till it gets to the destination network
 - From there, the destination network forwards the packet to the intended host

** Simplified model but sufficient for our purposes*

IP VERSION 4 (IPV4)

11000000	.	10101000	.	00000001	.	00001010	(Binary)
192	.	168	.	1	.	10	(Decimal)

- 32-bits
- Usually represented as four 8-bit bytes (octets)
- E.g. 206.109.1.6, 127.0.0.1, 192.168.1.10

IP VERSION 6 (IPV6)

An IPv6 address

(in hexadecimal)

2001:0DB8:AC10:FE01:0000:0000:0000:0000



2001:0DB8:AC10:FE01::

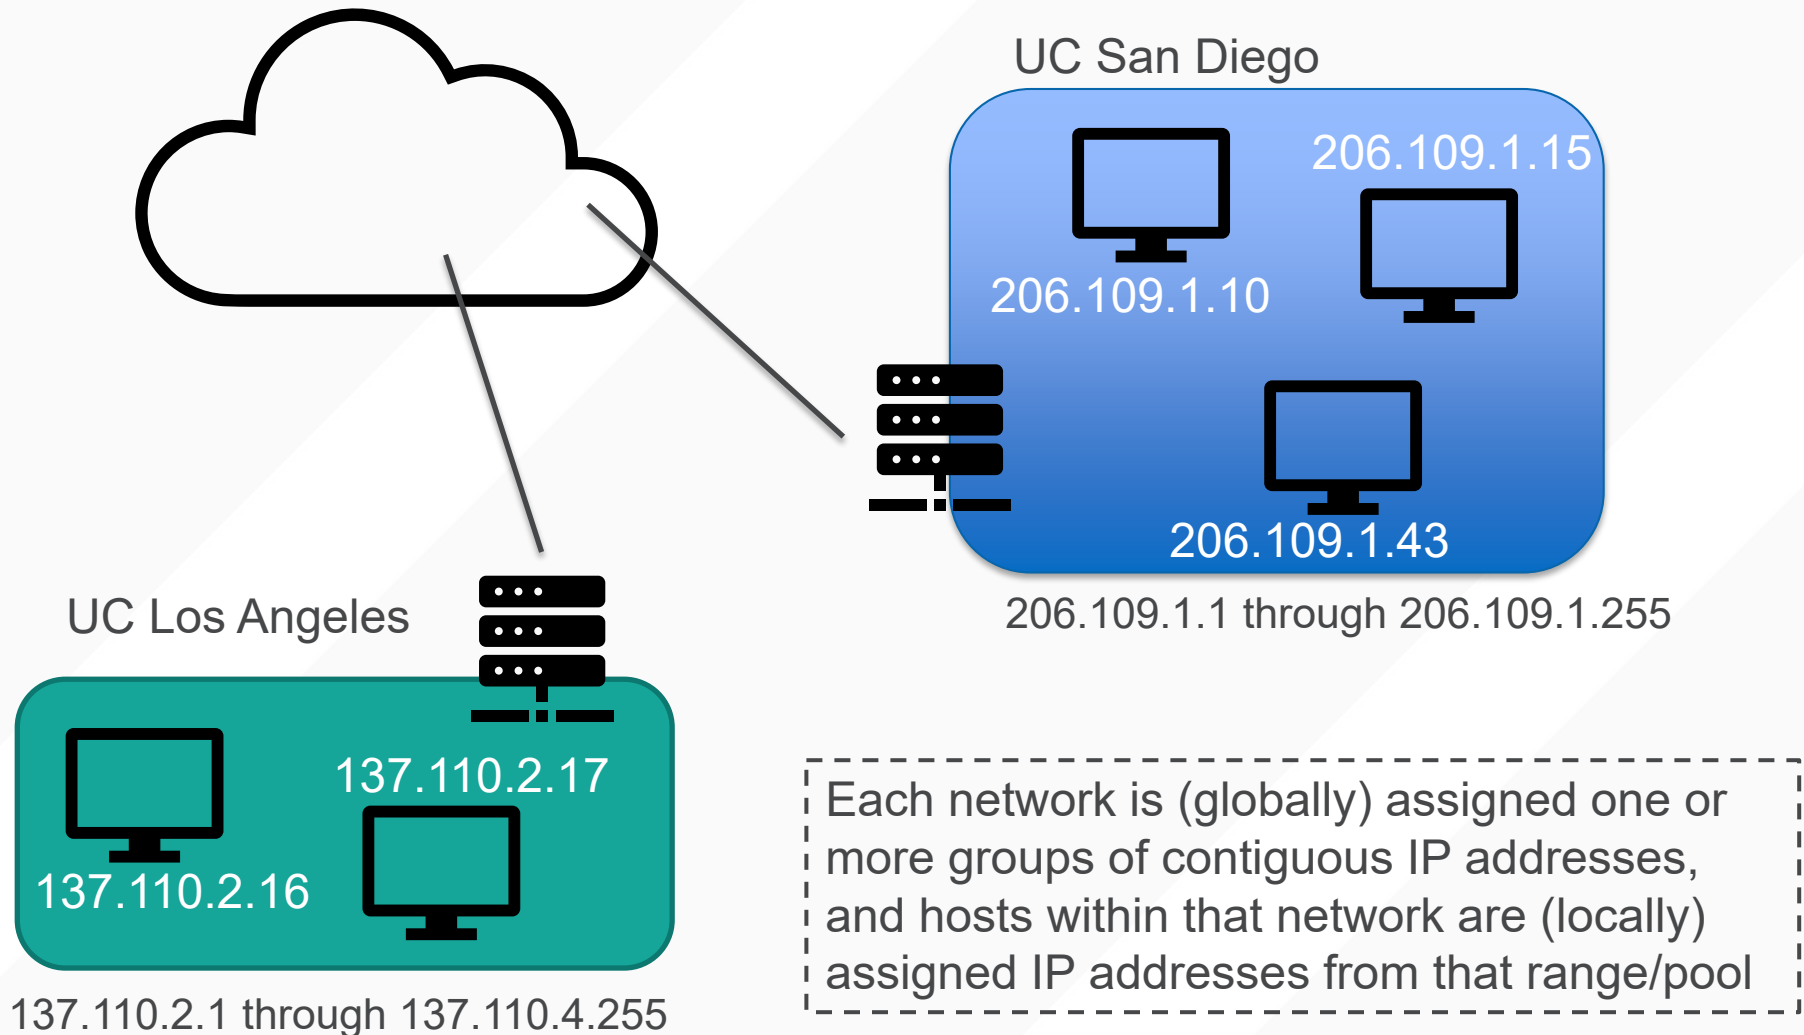
Zeroes can be omitted

Diagram showing the binary expansion of the IPv6 address. Four curved arrows point from the first four blocks of the compressed address (2001, 0DB8, AC10, FE01) to the first four blocks of the binary expansion. The binary expansion is shown in two lines: the first line contains the first 64 bits (0010000000000001:0000110110111000:1010110000010000:1111111000000001:) and the second line contains the remaining 64 bits (0000000000000000:0000000000000000:0000000000000000:0000000000000000).

0010000000000001:0000110110111000:1010110000010000:1111111000000001:
0000000000000000:0000000000000000:0000000000000000:0000000000000000

- 128-bits
- Represented as 8 16-bit blocks, in hex, separated by colons
- E.g. 2001:4860:4860::8888, ::1, 2345:0425:2CA1:0000:0000:0567:5673:23b5

THE INTERNET: A NETWORK OF SUB-NETWORKS



ROUTE AGGREGATION

- 2^{32} possible (IPv4) addresses spread across ~100,000 independent networks
- Each router keeps the “next hop” on a per-network basis, not per-host basis
- But networks can be different sizes (e.g. UCSD is bigger than a small startup)
- Each router has to keep a list of networks (and their next hops) + how “big” each network is

CLASS-BASED ADDRESSING (NOT REALLY USED ANYMORE)

- Most significant bits determines “class” of address

Class A

0	Network	Host
---	---------	------

127 nets, 16M hosts

Class B

		14	16
1	0	Network	Host

16K nets, 64K hosts

Class C

			21	8
1	1	0	Network	Host

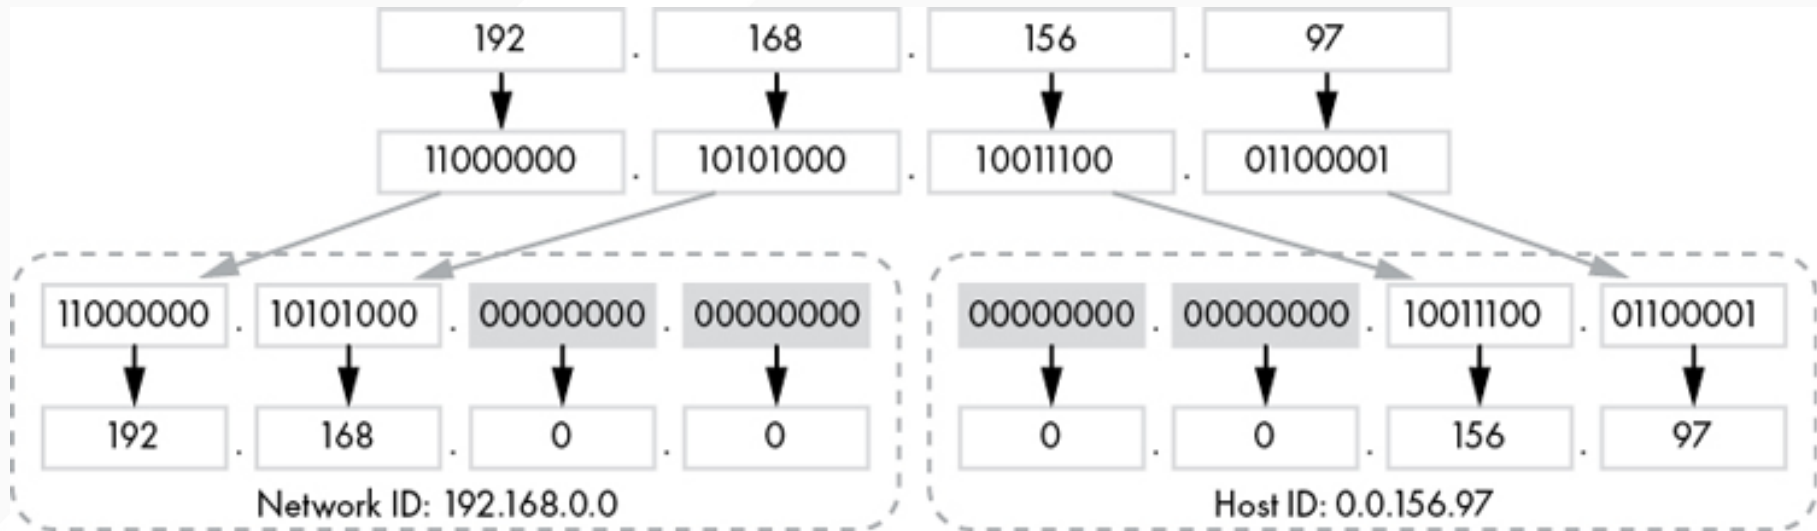
2M nets, 254 hosts

- Special addresses
 - Class D (1110) for multicast, Class E (1111) experimental
 - 127.0.0.1: local host (a.k.a. the loopback address)
 - Host bits all set to 0: network address
 - Host bits all set to 1: broadcast address

CLASS-BASED ADDRESSING EXAMPLES

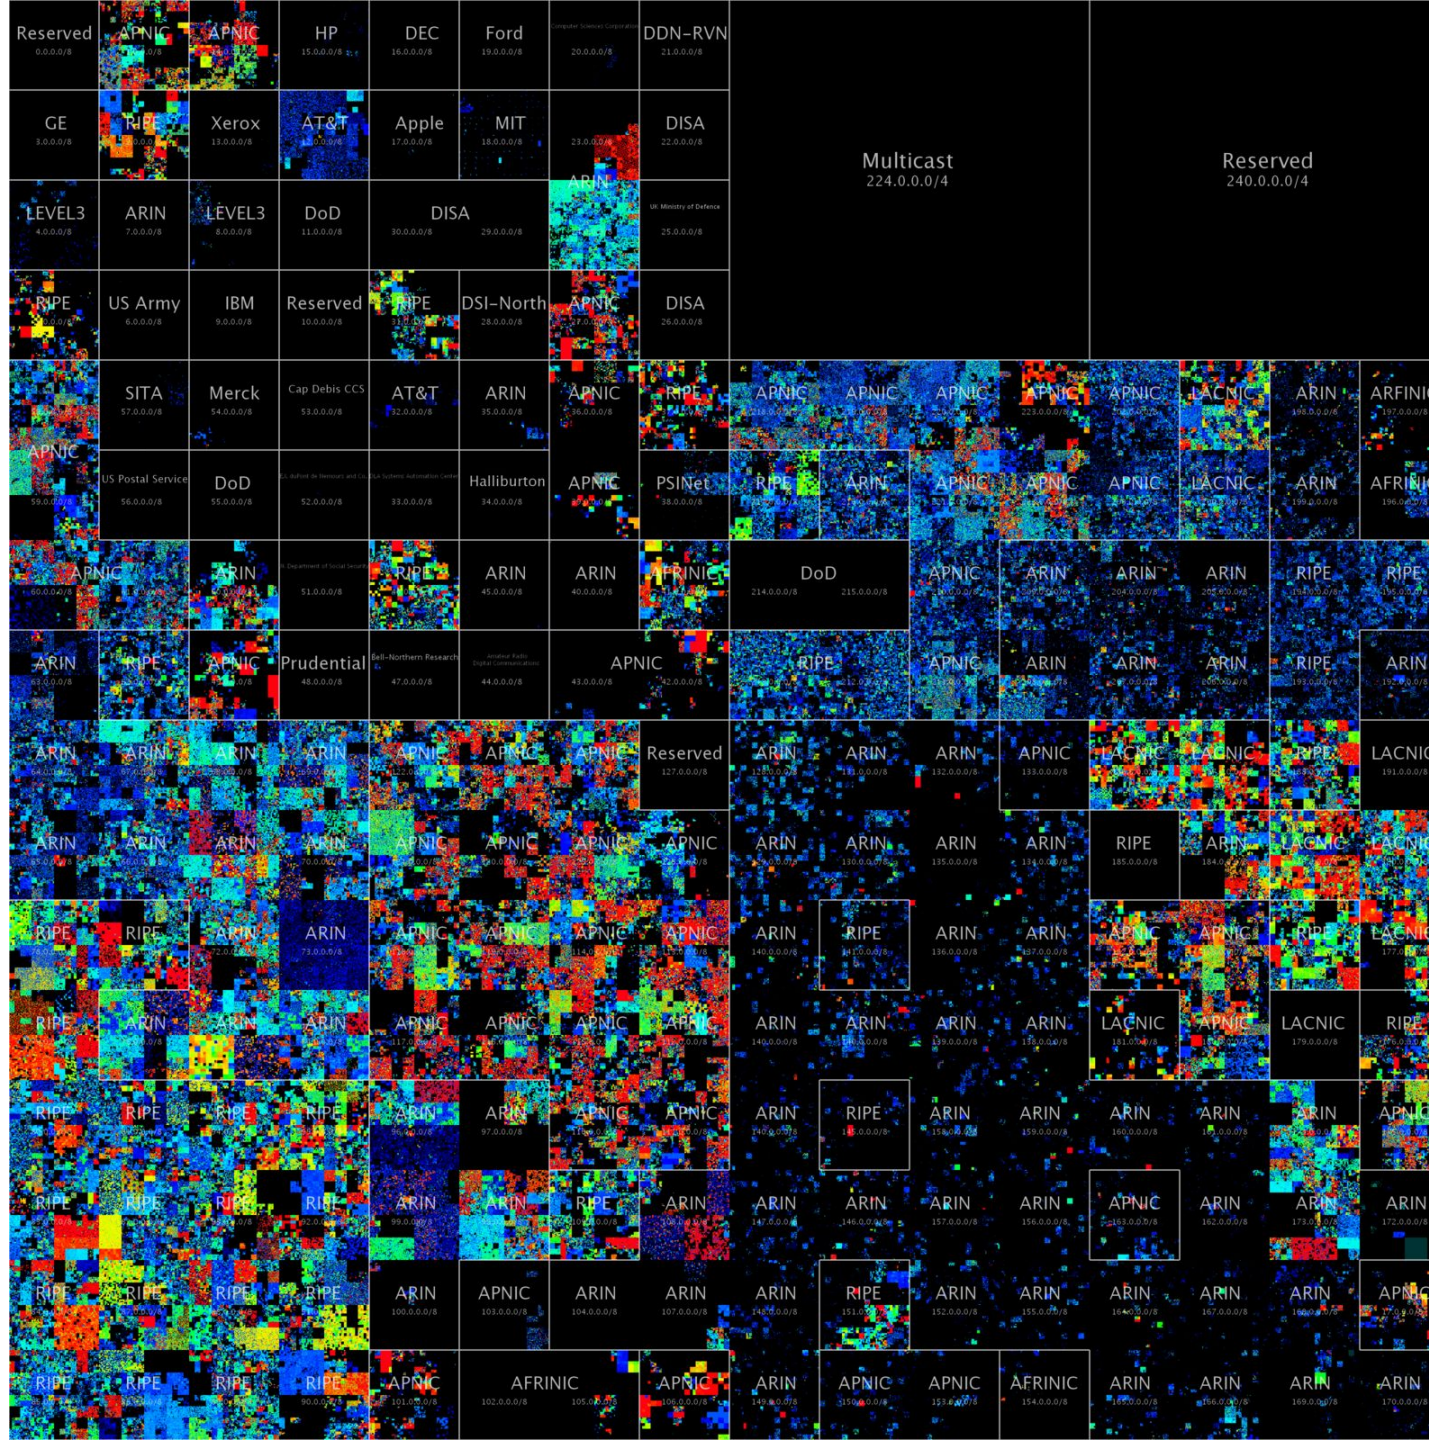
Network ID	First octet	Second octet	Third octet	Fourth octet	Host ID
8 bits	<div>Network</div> <div>10</div>	<div>Host</div> <div>1</div>	<div>Host</div> <div>2</div>	<div>Host</div> <div>3</div>	24 bits
16 bits	<div>Network</div> <div>172</div>	<div>Network</div> <div>16</div>	<div>Host</div> <div>1</div>	<div>Host</div> <div>2</div>	16 bits
24 bits	<div>Network</div> <div>192</div>	<div>Network</div> <div>168</div>	<div>Network</div> <div>1</div>	<div>Host</div> <div>2</div>	8 bits

ADDRESSING EXAMPLE



IP ADDRESS PROBLEM (1991)


- Address space depletion
 - In danger of running out of classes A and B
- Why?
 - Class C too small for most organizations (only ~250 addresses)
 - Very few class A – very careful about giving them out (who has 16M hosts anyway?)
 - Class B – greatest problem



IPv4 Census Map

June - October 2012

Utilization



Prefix Sizes



420 Million hosts that responded to ICMP Ping at least 2 times between June and October 2012
Source: Carna Botnet

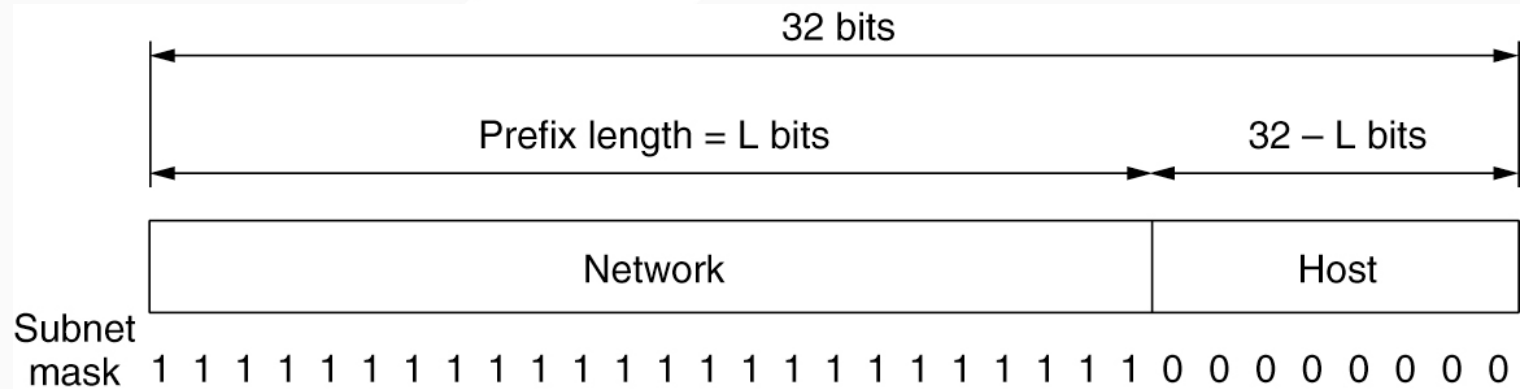
CIDR

- Classless Inter-Domain Routing (1993)
 - Networks described by variable-length prefix and length
 - Allows arbitrary allocation between network and host address

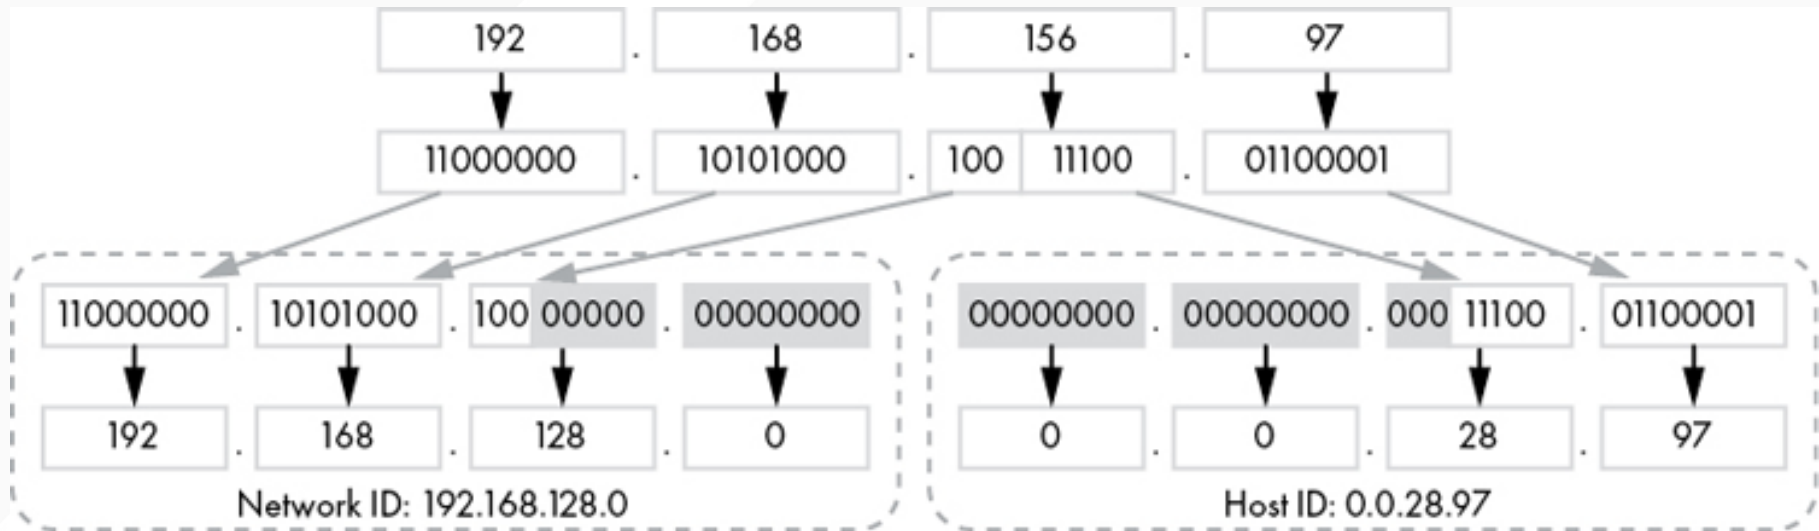


- e.g. 10.95.1.2 contained within 10.0.0.0/8:
 - 10.0.0.0 is network and remainder (95.1.2) is host
- Pro: Finer grained allocation; aggregation
- Con: More expensive lookup: **longest prefix match**

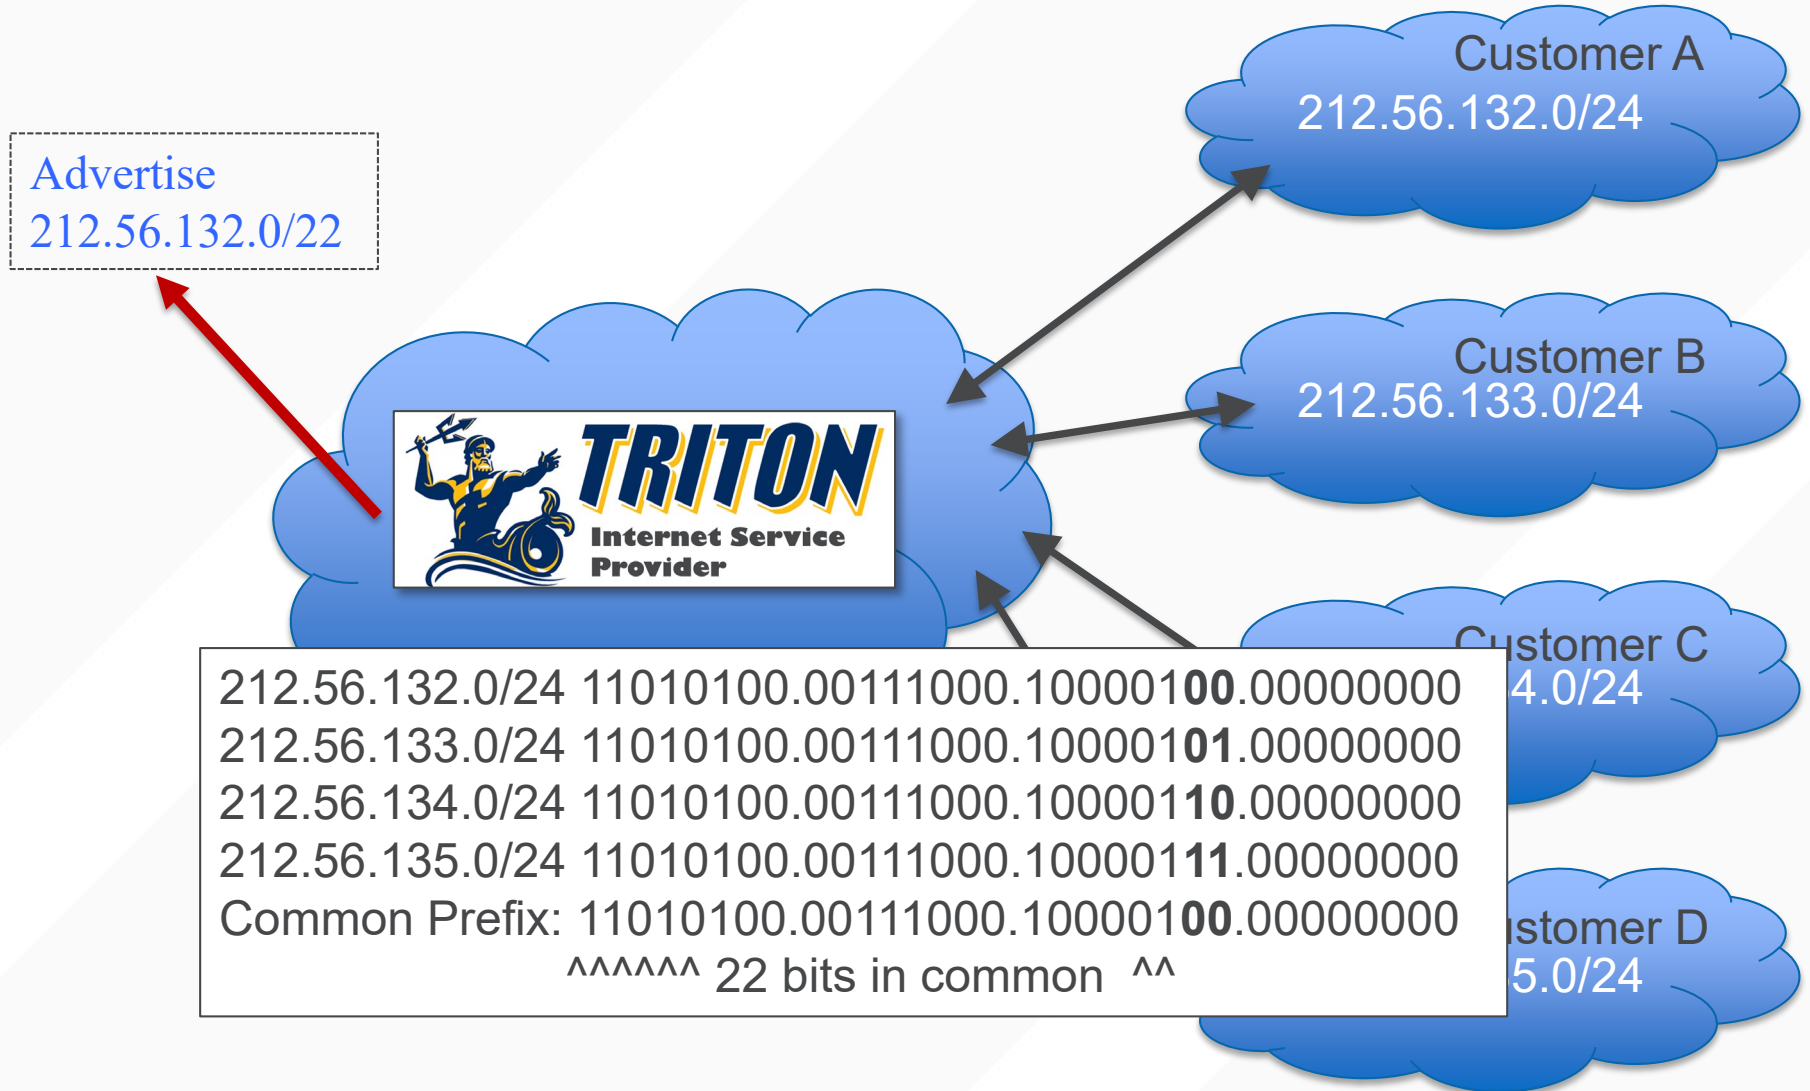
SUBNETS AND NETMASKS



192.168.156.97/19




ADDRESS AGGREGATION EXAMPLE



ROUTING TABLE EXAMPLE (R2)

- Packet to 10.1.1.6
- Matches 10.1.0.0/23

Destination	Next Hop
127.0.0.1	loopback
Default or 0/0	10.1.16.1
10.1.8.0/24	interface1
10.1.2.0/23	interface2
10.1.0.0/23	10.1.2.2
10.1.16.0/24	interface3




ROUTING TABLE EXAMPLE 2 (R1)

- Packet to 10.1.1.6
- Matches 10.1.1.4/30
 - Longest prefix match

Routing table at R1

Destination	Next Hop
127.0.0.1	loopback
Default or 0/0	10.1.2.1
10.1.0.0/24	interface1
10.1.1.0/24	interface2
10.1.2.0/23	interface3
10.1.1.4/30	10.1.1.101




ROUTING TABLE EXAMPLE 3 (H1)

- Packet to 10.1.1.6
- Direct route
 - Longest prefix match

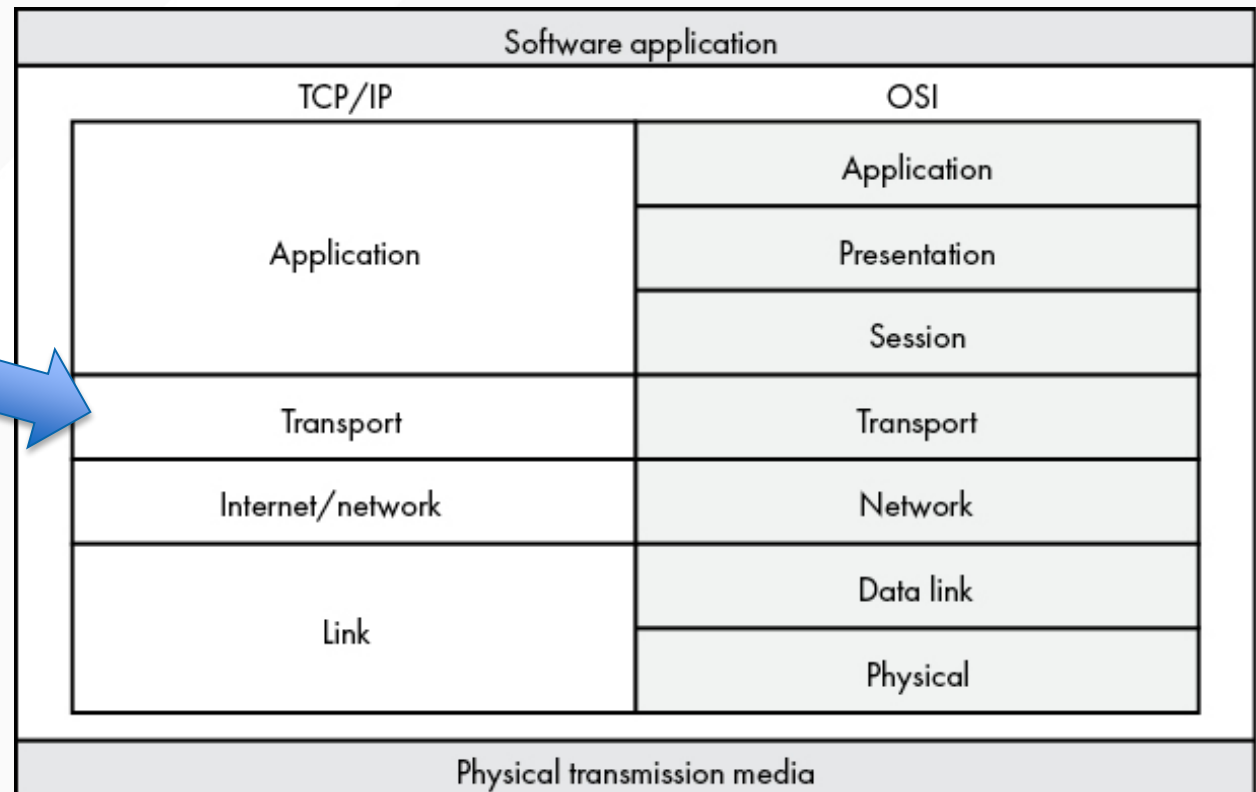

Routing table at H1

Destination	Next Hop
127.0.0.1	loopback
Default or 0/0	10.1.1.1
10.1.1.0/24	interface1
10.1.1.4/30	interface2



TCP/IP MODEL (VS OSI MODEL)

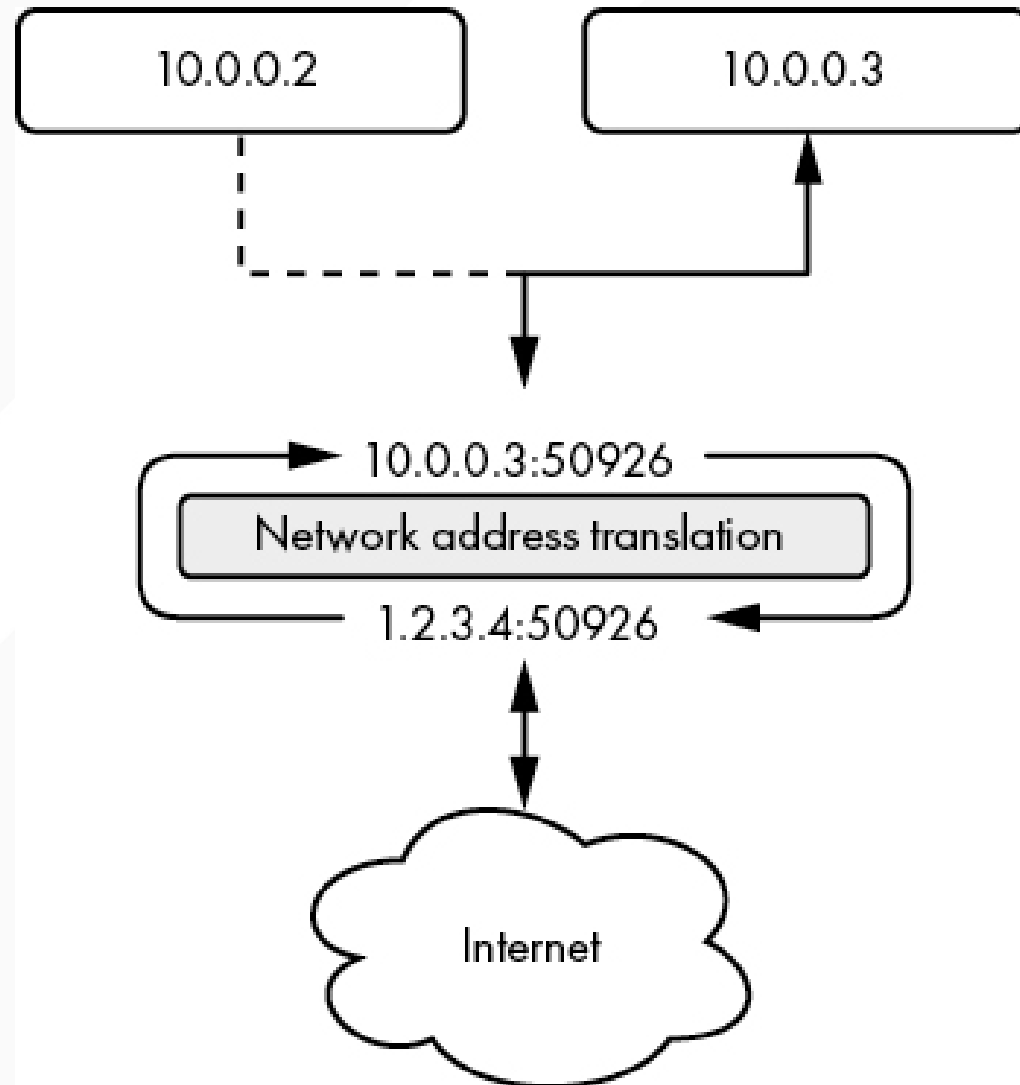
Which *process* on the destination machine?
What if I want to send a stream of bytes, not just a limited-size message?



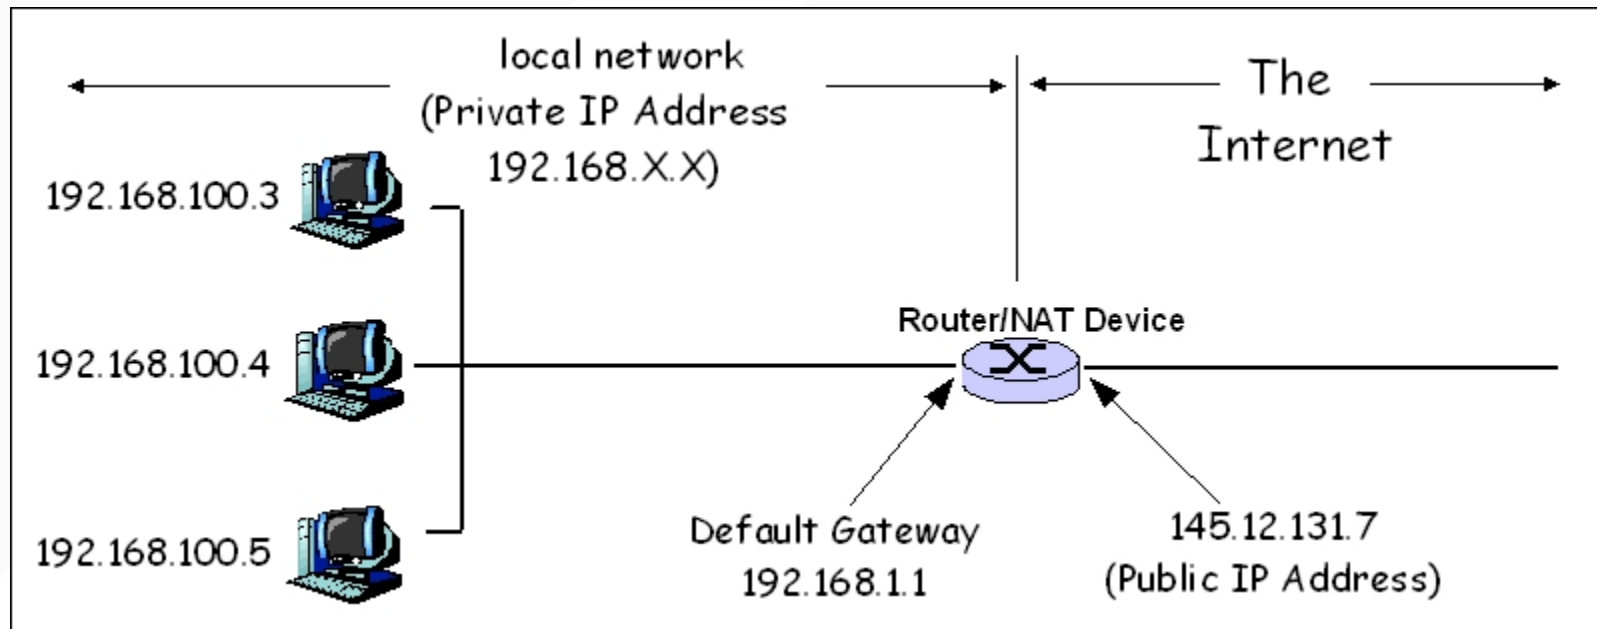
PORTS

- IP addresses identify a *machine*
 - Actually they identify a network interface on a machine
- How to identify different programs on the machine?
 - Process ID/PID? (no... why not?)
 - Instead we use a port (which is a 16-bit number)
 - 0-1024 reserved for the OS, you can use 1025-65535

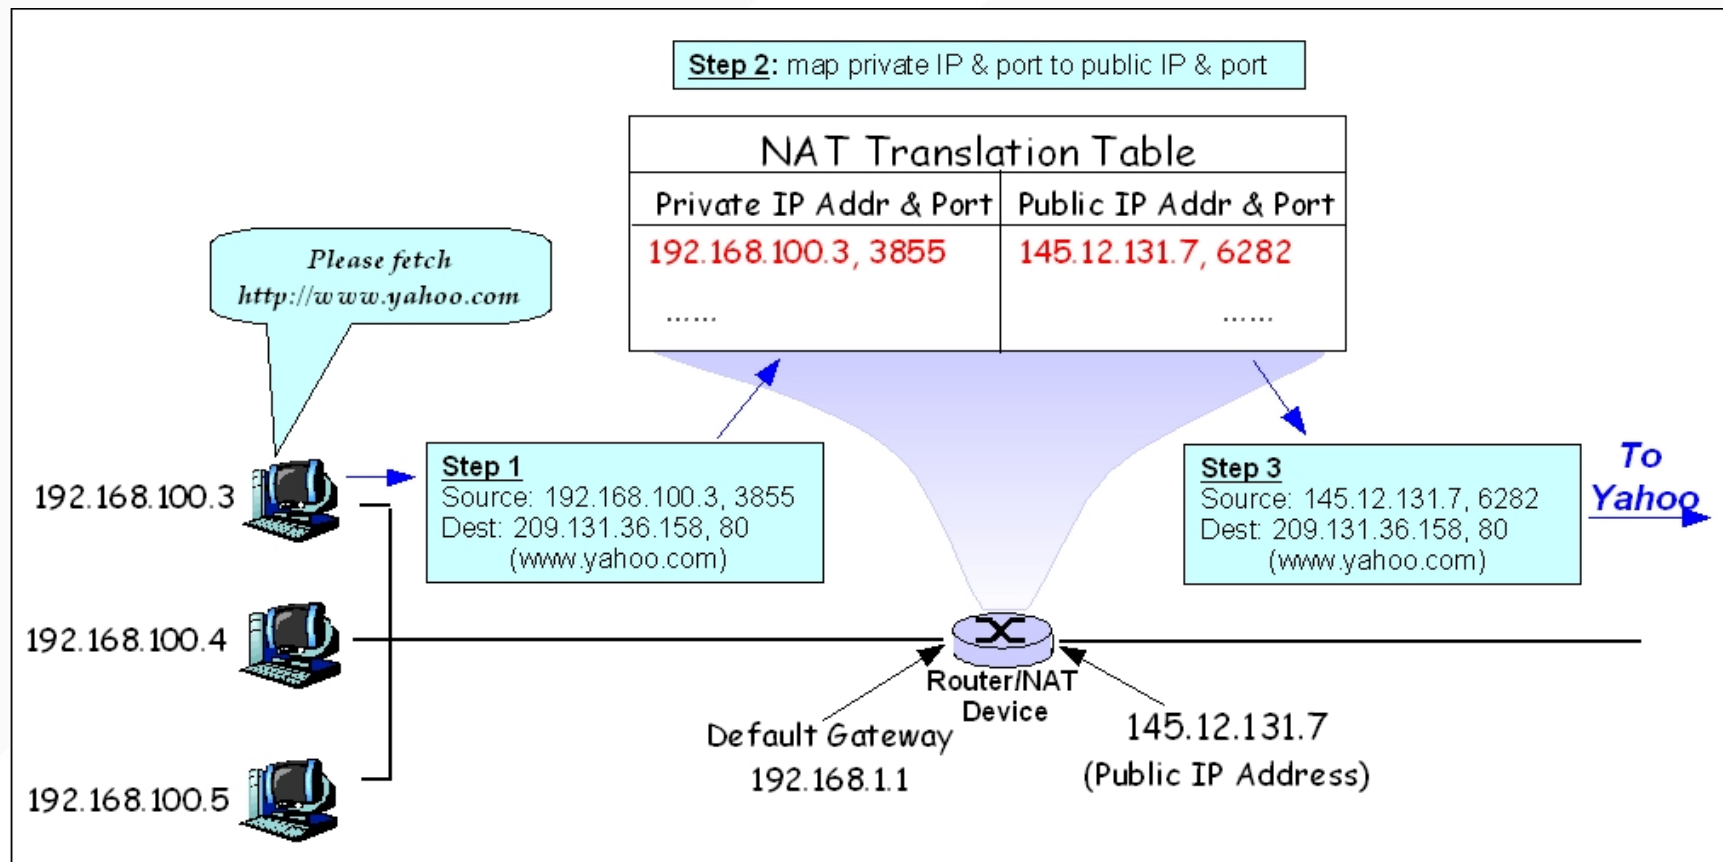
NETWORK ADDRESS TRANSLATION



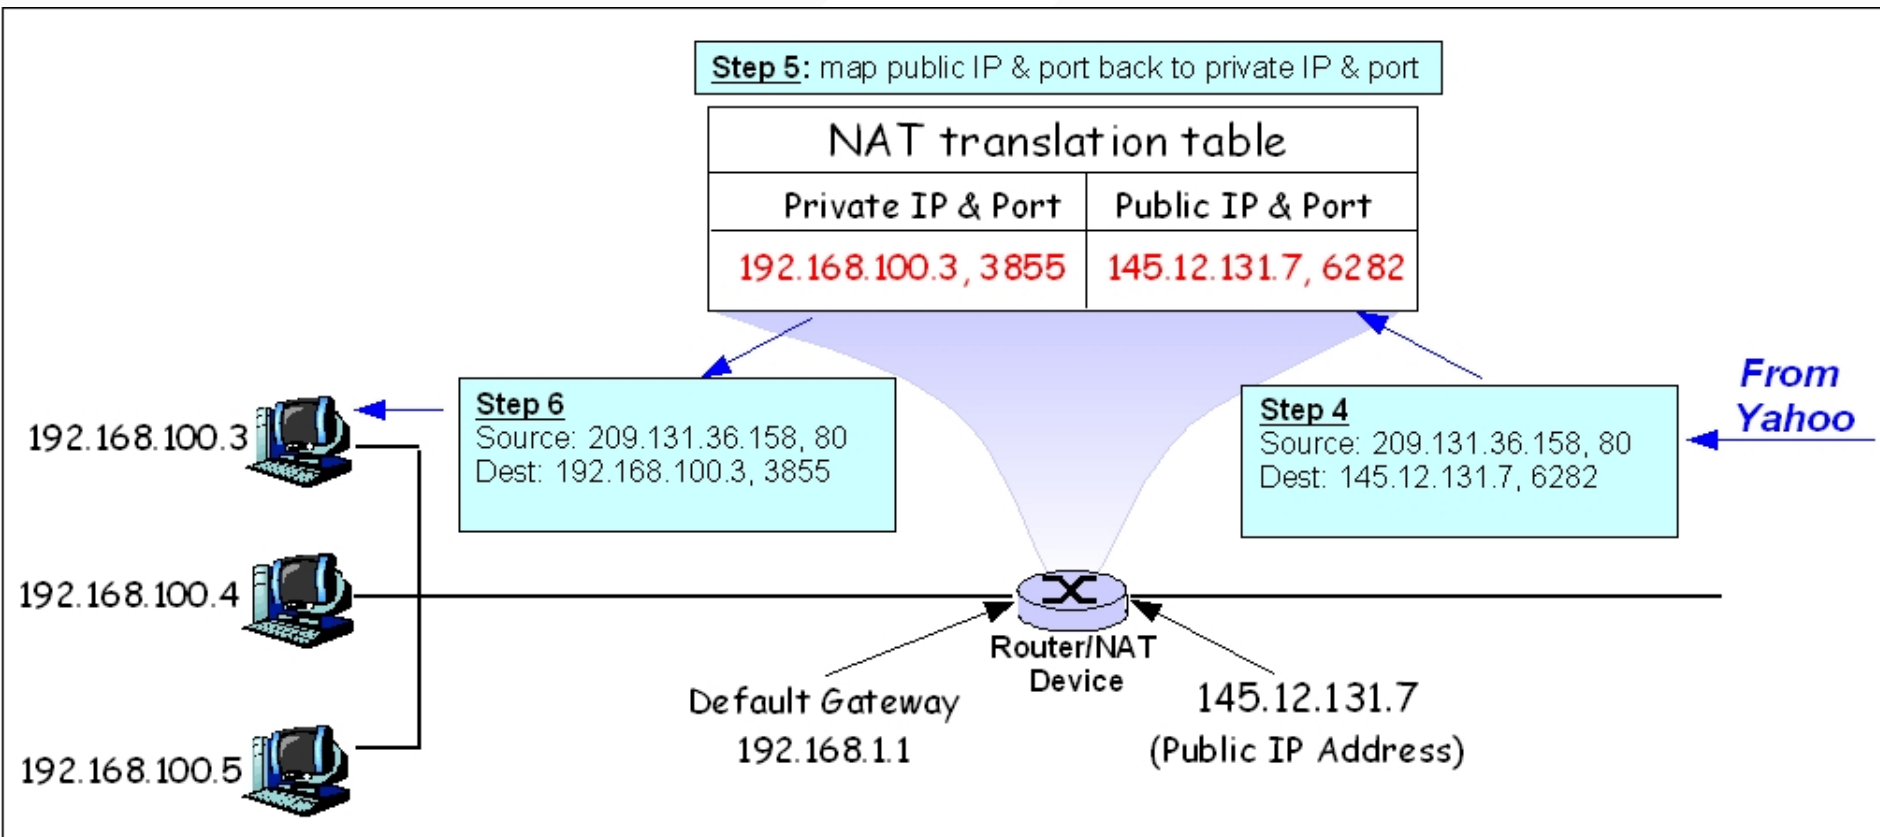
LOAD BALANCING VIA NETWORK ADDRESS TRANSLATION



LOAD BALANCING VIA NETWORK ADDRESS TRANSLATION



LOAD BALANCING VIA NETWORK ADDRESS TRANSLATION



UC San Diego